

**Flattering Advice: Avoiding Disappointment as a Driver of Gender
Discrimination**

Amanda Chen¹ & David Hagmann¹

¹ The Hong Kong University of Science and Technology

30 September, 2024

Author Note

Correspondence concerning this article should be addressed to Amanda Chen,
Department of Management, The Hong Kong University of Science and Technology, Hong
Kong. E-mail: zchengj@connect.ust.hk

Abstract

The purpose of advice is typically understood as helping recipients make better decisions and solve problems. However, in this paper, we propose that advisors also consider the recipient's belief utility and may try to avoid disappointing them, leading to what we term "flattering advice." Across three pre-registered experiments involving real interactions between advisors and advisees ($n = 3,900$), we demonstrate that individuals factor in others' expectations when giving advice, even when those expectations are not relevant to solving the decision-making problem. Specifically, advisors take advisees' expectations into account, tailoring their advice to create a more positive impression of the advisee and avoid causing disappointment. (Study 1). Advisors inflate their advice due to interpersonal concerns, with those who care about being liked by the advisee inflating their advice to a greater extent. Advisees, in turn, reward this flattery, perceiving these advisors as more likable and equally trustworthy compared to those who give more accurate advice (Study 2). Additionally, when advisors match advisees' expectations, men tend to receive more aspirational advice than women, due to their higher confidence, even when only expectations (not gender) are visible to the advisor. However, relying on these expectations leads to poorer decision outcomes (Study 3). We conclude by discussing the practical and theoretical implications of our findings, as well as directions for future research.

Keywords: Gender difference, Advice, Interpersonal relationship

Word count: 4,362

Flattering Advice: Avoiding Disappointment as a Driver of Gender Discrimination

Introduction

Seeking and giving advice can present conversational challenges that are difficult to navigate. Advice seekers may be embarrassed to ask for help (Bénabou et al., 2022), and they may fail to do so if the information they receive could lead to disappointment (Gill & Prowse, 2012). Indeed, people prefer advisers who withhold unpleasant information (Shalvi et al., 2019) and punish those who tell them things they do not want to hear (John et al., 2019). People appear to engage in strategic selection of advice to avoid unpleasant information and protect the utility they get from positive beliefs about themselves and the future they anticipate (Loewenstein, 1987; Loewenstein & Molnar, 2018). But do advice givers anticipate and consider the hedonic cost of honest information to the recipient? While previous work has emphasized the instrumental value of advice (e.g., satisfaction with a decision or accuracy of a prediction), we focus on the motivations of the person providing advice.

Research on advice has largely focused on settings that do not tackle belief utility. For example, participants may have been asked to estimate the number of balls in an urn or choose from pairs of lotteries (Benjamin & Budescu, 2015). Advisers have been shown to be motivated to give accurate advice (Jonas & Frey, 2003), and engage in perspective-taking to meet the preferences of the advisee (L. J. Kray, 2000; L. Kray & Gonzalez, 1999). When they are uncertain, they even prefer that advisees do not rely (extensively) on their recommendations (Ache et al., 2020). Thus, it seems that people are concerned about providing good advice even when they are not incentivized for the recipient's decision quality.

In many real settings, however, advice incorporates beliefs about the recipient. Someone recommending an expensive restaurant conveys that they think the advisee is wealthy enough to afford an expensive meal, and someone recommending a student pursue a PhD in economics assumes that they are intellectually capable. Conversely, telling a poorly

performing student to switch to a less-demanding major conveys a belief that the student is not capable of improving. Thus, while the painful truth may prevent someone from making a bad decision, it may come at a hedonic cost to the recipient and damage the relationship between the adviser and advisee. We propose that advisers take these costs into account, even when the interactions are entirely anonymous and they are incentivized for the quality of the advisee’s decision.

Specifically, we focus on the signaling value of advice: what does the advice imply about the adviser’s perception of the advisee? We propose that advisees have expectations about the advice they receive and treat this as a reference point (Kőszegi & Rabin, 2006). Advice that falls short of these expectations creates disappointment, while advice that exceeds them is flattering and creates positive interpersonal benefits because it creates a perception that the adviser holds the advisee in high esteem. Such interpersonal considerations may be as important, if not more important, to the adviser than the instrumental value of the advice: the adviser directly experiences a souring relationship, but may suffer little and only much delayed consequences if an advisee makes a poor life decision.

Our account therefore predicts that people give and receive advice that is more flattering and suggestive of inflated ability and optimistic outcomes than would advice in the absence of such interpersonal considerations. However, we propose that this is not true for all advisees equally. Presenting flattering advice requires the formation of beliefs about the advisee’s expectations (reference point). Exley and Nielsen (2022) show that men are more overconfident than women, conditional on equal performance, and that people anticipate this difference. An adviser who seeks to avoid disappointment and incorporates expectations into their advice would then present more favorable advice to male advisees than to female advisees (see Figure 1 for an illustration). This provides a novel account for a well-known finding that men receive more aspirational advice than do women (Kanze et al., 2018). Notably, when it pays to have accurate beliefs, such flattering advice would come at a cost:

while men would be more likely to aim higher, they would also suffer greater costs from aiming too high and failing to meet the objective.

In a large, preregistered experiment ($n = 1,200$), we show that advisers consider the expectations of advisees, even when such expectations do not provide informative insights for the underlying decision-making problem. We find that as a result, showing expectations (but not gender) to advisers leads men to receive more aspirational advice than women.

Incorporating expectations leads to worse advice, even when advisers are incentivized based on the outcome of the advisee's decision. In a second preregistered experiment ($n = 700$), we show that advisers respond to incentives for interpersonal considerations. They tell people that they are more attractive than they believe them to be, and inflate to a greater extent when incentivized for likability rather than accuracy. Notably, they inflate even when only incentivized for accuracy. Advisees indeed reward such flattering advice, rating the advisers more likeable and no less trustworthy than those who provide more accurate advice.

Open Science Statement

We report all manipulations, measures, and data exclusion in this and the following study. The preregistration reports, screenshots of all experimental materials, and the analysis code to replicate all statistical analyses and figures are available on the Open Science Framework (https://osf.io/8r3d4/?view_only=5ad7bafcd16b4d4ba08bb28b0e2bd02d).

Study 1

In Study 1, we investigate whether givers take into account advisees' expectations when providing advice. We compare the advice given to recipients who expressed either high or low expectations, across two treatments: one where givers were aware of these expectations (Expectation treatment) and one where this information was unknown to them (Performance treatment).

The experiment consisted of three stages, in which we manipulated advisees' expectations and whether givers had access to this information. First, we measured

participants’ performance, then informed them that their performance would be compared to that of another participant, with the higher scorer eligible to win a bonus. Participants were asked to choose their competitor from either the top 20 performers (“High Performer Group”) or the bottom 20 performers (“Low Performer Group”) out of a pool of 50 competitors. We nudged participants to have high or low expectations by encouraging them to compete against either the High Performer Group or the Low Performer Group. Advisors were randomly assigned to one of two treatments: seeing only performance information or both performance and expectation information. We hypothesize that advisees who indicate they are capable of competing against strong competitors will be more likely to receive advice to do so.

Methods

As part of our main experiment, participants were asked to choose a group of previous participants to compete against. To create these groups, we first recruited a sample of 50 participants from Prolific to complete a 10-question ability quiz. The quiz included word puzzles, identifying emotions from photos, and selecting the best responses in given scenarios. These questions were adapted from surveys that measure problem-solving, emotional intelligence, and communication skills. We selected these specific abilities based on a pretest in which participants ranked them as the three most critical skills in modern society, making performance on this task ego-relevant. Participants had five minutes to answer the quiz and were paid 5 cents for each correctly answered question. We define the top 20 scorers as the “High Performer Group” and the bottom 20 scorers as the “Low Performer Group.”

We then recruited 201 participants to serve as advisees in Stage 1 of our main experiment. These participants completed the same 10-item ability quiz as the earlier group, with the opportunity to earn 5 cents for each correct answer. After the quiz, we asked them to indicate their initial preference for competitors. Half of the participants were randomly shown the default option of the ‘High Performer Group,’ highlighted in blue (“High Default”

treatment), while the other half were shown the ‘Low Performer Group’ as the default (“Low Default” treatment). This manipulation was designed to influence advisees’ expectations, as advisees who believed they could compete against stronger competitors were perceived by givers to have higher expectations. However, participants were free to opt out of the default choice.

We then recruited 1,000 participants for Stage 2, placing them in the role of advisers. We began by informing them of the ability quiz that participants in the preliminary study and Stage 1 had completed, and informed them of the average score of all participants in the preliminary study. Advisers had to recommend whether an advisee should compete against the Low Performer Group or the High Performer Group (we used these terms in the survey). We anticipated that being told to compete against High Performers is more flattering and hence being told to compete against the Low Performers would be disappointing if one had expected to do well. Advisees would earn 20 cents if their score was equal to or higher than a randomly selected member of the Low Performer Group and they chose to compete against this group, or 50 cents if they performed as well or better than a randomly selected member of the High Performer Group and they chose to compete against that group. If they scored lower, they would not receive any bonus.

Advisers were randomly assigned to one of two treatments. In the “Performance” treatment, they only observed the score of the advisee on the ability quiz. In the “Expectation” treatment, they observed the score as well as the advisee’s initial preference for which group to compete against. However, since the outcome was based solely on the quiz score and advisees didn’t know their scores, their initial choice wasn’t useful in deciding which group had a higher expected bonus. Advisers gave recommendations to 10 advisees and were told that if their advice was used by a participant in a follow-up survey, they would receive the same bonus as that participant. The survey concluded with basic demographic questions.

Finally, we invited participants from Stage 1 back for the follow-up survey. Following our preregistration, we kept the survey open for 7 days. In total, 176 participants (86 men, 90 women) returned. The brief survey reminded them of the task they completed in Stage 1, informed them that other participants from Prolific had observed their real score and given them advice against which group to compete against, and finally were reminded them of how many questions they guessed they had answered correctly. Importantly, they were not informed of their true score or the score of the groups they could compete against. Participants then observed the advice from a randomly selected adviser and made their decision.

Results

We begin by examining the performance of the Stage 1 participants. Because our treatment takes place after participants completed the ability quiz, we would not expect a difference in performance across the Low and High Default treatments. Indeed, the two groups scored no different from one another (5.10 and 5.18 for the Low Default and High Default treatments, respectively, $t(199) = 0.31$, $p = .756$). The priming did, however, affect advisees' initial choices. In the "Low Default" treatment, only 22% of participants initially preferred competing against the High Performer Group, compared to 65% of participants who saw the High Performer Group highlighted ($\chi^2(1, n = 201) = 36.62$, $p < .001$).


Next, we examine whether advisers took the advisees' expectations into account. Recall that advisees who defaulted to competing against the High Performer Group were indeed more likely to choose this group, which, from the advisers' perspective, could signal high confidence in their abilities. Our theory predicts that showing expectations should lead to more flattering advice—specifically, a recommendation to compete against the High Performer Group—for advisees with a high default choice compared to those with a low default choice. Column 1 of Table 1 reports a linear probability model with advice to compete against the high group as the outcome measure, incorporating the default choice

Table 1

Column 1 displays the advice given to compete against a High Performer Group based on whether the advisees were primed with low or high default and the initial preference shown to the advisor. Column 2 displays the advice given to compete against a High Performer Group based on the default and advisees' performance when their initial preference was visible to the advisor. All standard errors are clustered at the level of the advisor.

	(1)	(2)
High Default	0.011 (0.013)	0.025* (0.011)
Expectation Shown	-0.041* (0.016)	
Expectation x High Default	0.037+ (0.019)	
Performance		0.168*** (0.004)
Constant	0.389*** (0.012)	-0.503*** (0.020)
N	10 000	5000
+ p < 0.1, * p < 0.05, ** p < 0.01, *** p < 0.001		

advisees saw, whether expectations were shown to the adviser, and the interaction of the two. Because each adviser made ten recommendations, we cluster standard errors at the adviser level. As predicted, we find a marginally significant interaction effect: advisees with a default choice to compete against the High Performer Group (nudged to have high expectations, though not necessarily holding them) were more likely to receive advice to compete against stronger competitors when their expectations were shown compared to when they were hidden. We reported a linear probability model on advising to compete against the High Performer group for advisers in the “Expectations” treatment, controlling for the true performance of the advisee in Column 2 of Table 1. We observe a significant main effect of the default treatment, where a high default increased the likelihood of receiving flattering advice. It is worth noting that, since only 65 % of advisees in the high default group actually stuck with their default choice, this represents a conservative test of the effect of high expectations on the likelihood of receiving flattering advice.

**Figure 1**

Advisees who were nudged to express high expectations were more likely to receive advice to compete against the High Performer Group.

Discussion

When the quality of a decision does not depend on someone's expectations, the adviser's awareness of the advisee's expectations should not influence the advice given. However, in line with our argument that advisers take into account the belief utility of the advisee and prefer to flatter rather than disappoint, we find that displaying the advisee's expectations does make a difference. Specifically, advisees who expressed a preference for stronger competitors were more likely to receive 'flattering' advice, encouraging them to pursue that option. One possible explanation for this finding is that, rather than trying to flatter, advisers may simply be telling advisees what they want to hear. In the next study, we aim to investigate the mechanism behind this discrepancy in advice-giving when expectations are revealed, exploring the question: Why do people give flattering advice?

Study 2

So far, we have shown that displaying expectations influences the advice given. Specifically, advisers were more likely to recommend competing against a group of high

performers to advisees with higher performance expectations compared to those without. We now directly test whether this effect is driven by interpersonal concerns. In particular, we explore whether more flattering advice stems from a desire to be liked, and whether advisees, in turn, like advisers more when they receive flattering advice.

We conducted a three-stage experiment in which we experimentally manipulated the incentives for advisers. We began by inviting a group of advisees to upload photos of themselves (“selfies”) and informed them that they would be rated on their attractiveness. We grouped them with nine other participants of the same gender and recruited participants of the opposing gender to rank them from most to least attractive and to provide advice. Specifically, we asked them to advise the participant they ranked as the 7th most attractive (4th least attractive) on what rank they should bet they were ranked by a larger group of raters. Advisers were randomly assigned to two treatments, receiving a bonus payment either if the advisee guessed their rank accurately or if the advisee evaluated the adviser as likeable as measured by an unincentivized scale response. We hypothesize that advisers who want to be liked by advisees will recommend that they bet on a lower rank, i.e. that they are more attractive.

Methods

We recruited 300 participants from Prolific and, after asking demographic questions, invited them to upload photos of themselves (selfies) to be rated by other participants on attractiveness. 100 men and 107 women agreed to do so and uploaded pictures that adhered to our instructions (e.g., did not include other people). We selected the first 100 photos from women to arrive at a gender-balanced sample ($M_{\text{Age}} = 39.37$ years). Participants were informed that their selfies would be randomly grouped with those of nine other participants of their gender and ranked in terms of attractiveness by a group of new Prolific participants of the opposite gender. Lastly, they guessed their rank (unincentivized).

Next, we recruited a new, gender-balanced sample for the role of advisers. 472

participants from Prolific ($M_{\text{Age}} = 41.03$ years; 49.79% Female) started by providing demographic information, then were matched to a group of the opposite gender. They then ranked participants from most to least attractive using a drop-down menu next to each picture. Because of a limitation with the survey software, participants were able to select the same rank for multiple participants, rather than rank them uniquely. We remove 115 participants who did not follow instructions and provide a unique ranking.

After submitting their ratings, they saw the photo of the participant they had ranked as the 7th most attractive (i.e., the 4th least attractive). They were reminded of the rank they had given to that person and informed that this participant would be invited back and could earn a \$1 bonus if they guessed their rank correctly. The rank was determined by the aggregate ratings of all participants who had ranked this group. Their task was to give advice to this participant about the rank they should bet on based on having observed all ten selfies and their own assessment. We randomly assigned participants to one of two incentivization schemes. In the “Accuracy” treatment, they received a bonus identical to the advisee: \$1 if they guessed their rank correctly. In the “Likeability” treatment, we informed them that they would be rated by the advisee on a 5-point Likert scale for how likeable they thought they were. Each point on the scale would translate to a bonus of 20 cents. They then selected a rank that they would recommend the advisee to bet on.

Finally, we invited participants from Stage 1 and were ranked as the 7th attractive by at least one adviser (so that they received advice) back for the follow-up survey. Following our preregistration, we kept the survey open for 7 days. In total, 146 participants (77 men, 69 women) returned. They were reminded of the selfie they uploaded in Stage 1 and informed that a group of 10 selfies, including theirs, had been rated by other participants from Prolific. They then saw the advice from a randomly selected adviser and made their estimate with a \$1 incentive for guessing accurately. They then saw the rank they had been advised to bet on one more time and were asked to rate the adviser’s likability, warmth,

friendliness, good-naturedness, trustworthiness, and sincerity on 5-point Likert scales (adapted from Fiske et al. (2007)).

Results




Figure 2

Advisers were asked to suggest what rank to bet on for someone they ranked as the 7th most attractive out of ten people. When incentivized for accuracy, most of them advised betting on 7, which was based on their evaluation of the recipient's attractiveness. However, when incentivized for likability, many participants gave flattering advice by suggesting betting on higher ranks that imply a higher evaluation of attractiveness.

We begin by examining the prior beliefs of advisees who uploaded their selfies. On average, men guessed they ranked 5.76 in their group of 10 and women guessed that they ranked 6.37. Participants overall underestimate their attractiveness relative to the benchmark average of 5.5. Moreover, women do so more than men $t(205) = 2.15, p = .033$. Notably, people's self-perceptions correlated strongly with the aggregate ratings of the advisers ($0.43, t(193) = 6.70, p < .001$). However, there was substantial heterogeneity in perceptions of attractiveness. Of the 200 participants, 149 were ranked as 7th most attractive by at least one adviser. On average, men in this subset estimated they were ranked 5.94th and women estimated they were ranked 6.59th ($t(143) = 1.94, p = .054$).

Next, we turn our attention to the advisers (Figure 2). In the Accuracy condition, those uploading selfies were advised to bet on rank 6.19. Notably, this is significantly more attractive than the 7th rank those advisers had themselves guessed just on the prior screen ($t(234) = -8.82, p < .001$). This suggests that even when incentivized for accuracy, participants offered flattering advice.¹ Importantly, and as predicted, we find that advisers in the Likeable treatment recommend betting on a lower rank, communicating that they think the participant in the selfie is more attractive ($5.38, t(470) = 5.44, p < .001$). This suggests that advisers inferred that advising someone that they were more attractive would make the adviser appear more likeable and therefore provided advice that communicated a more favorable impression of the participants' attractiveness. The distribution shown in Figure 2 shows that participants do not simply tell participants that they are the most attractive person in the group. They may infer that flattering advice needs to be somewhat realistic to be believable. We return to this in the general discussion.

Table 2

When individuals receive advice that implies a high level of attractiveness (lower rank), they tend to perceive the advice giver as more likable (Column 1), warm (Column 2). Column 3 shows that advisors are rated as more trustworthy when they advice lower ranks, but this relationship is only directional.

	(1)	(2)	(3)
Advised Rank	-0.109* (0.046)	-0.133** (0.044)	-0.054 (0.042)
Constant	3.777*** (0.279)	3.903*** (0.266)	3.397*** (0.257)
N	146	146	146
+ $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$			

Finally, we examine whether flattering advice indeed leads to more positive evaluations of advisers, or whether flattering advice is dismissed as insincere. Following our preregistration, we average the ratings on likeability, warmth, friendliness, and

¹ This shading could be due to concerns of avoiding disappointment. However, it could also be that advisers are uncertain about the rankings they had given and make a recommendation that combines their own belief with a uniform prior.

good-naturedness to create a scale of likeability ($\alpha = 0.92$); and we create a scale of trustworthiness by averaging the ratings of trustworthiness and sincerity ($\alpha = 0.87$).




Figure 3

Participants tend to rate advisors who suggest betting on a higher rank (implying greater attractiveness) as more likable and warm.

As seen in Figure 3, we found that advisees who suggested that the advisee was more attractive (lower rank) were indeed rated as more likeable and warm ($b = -0.109$, $p < .05$; $b = -0.133$, $p < .001$, respectively; See Columns 1 and 2 of Table 2). When using their self-expectations as a reference point, advisors who suggest a rank lower than that point are perceived as less likable and warm (see Figure 4). Interestingly, these benefits are not at the cost of being hypocritical; advisers who recommend a more favorable rank are viewed as no less trustworthy (Column 3 of Table 2).

We were not powered to do a comparison across the two experimental groups and did not preregister such a difference. Indeed, we find no difference in likeability and warmth across the two treatments ($t(144) = 1.09$, $p = .275$, and $t(144) = 0.58$, $p = .562$, respectively). We also assess the quality of advice by measuring the discrepancy between the advised rank and the final rank. Our analysis revealed no significant difference between the two

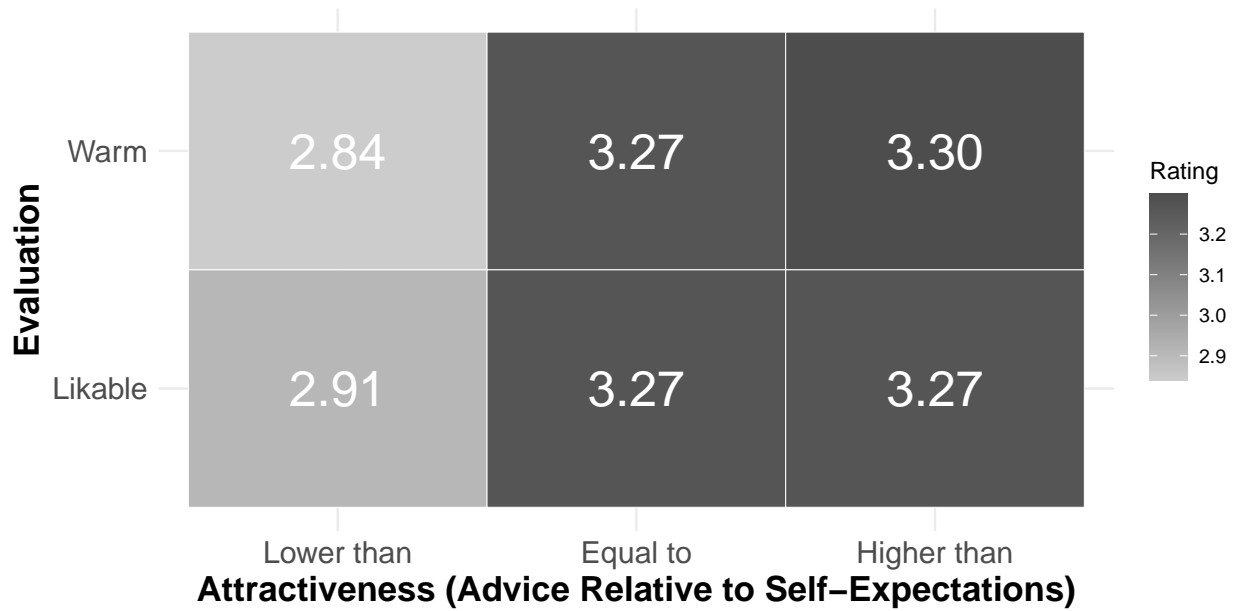


Figure 4

After participants received advice that suggested betting on a rank that was lower than, equal to, or higher than their self-expectations, they evaluated advisors based on their likability and warmth. Participants rated advisors who gave flattering advice and delivered a high evaluation of their attractiveness as more likable and warm than those whose advice implied a lower evaluation of their attractiveness.

treatments in terms of how helpful the advice was for advisees in making accurate guesses, $t(144) = -0.29, p = .776$. ## Discussion

In a context where advice communicates ego-relevant information (here, people's attractiveness), we find that the advice people give is contingent on their incentives. Specifically, when they get rewarded for being more likeable, they recommend that the advisee bet on a more favorable rank than when they are incentivized for accuracy. Importantly, advisees do not discount flattering advice and instead evaluate people who advise them to bet on a more attractive rank as warmer and more likeable. These gains to interpersonal perceptions do not come at the cost of trustworthiness, even when the advice is substantially inflated (e.g., a recommendation that the advisee is the most attractive in the group).

Study 3

We begin by examining whether advisers take the expectations of advisees into account, even when those expectations are not relevant to the decision the advisee has to make. In a three-stage experiment, we first recruit a sample of advisees and ask them to complete a mathematics quiz. After answering ten multiple-choice questions, participants are anchored to low or high performance expectations and asked to guess how many questions they answered correctly. Next, we recruit advisers and show them either only the true performance of the advisee, or the true performance and how many questions they guessed they answered correctly. Advisers are tasked with recommending whether the advisee should compete against a group of high-performers or a group of low-performers on the mathematics quiz, where a potential bonus depends on outperforming the competitor and which group was chosen. Finally, we invite the advisees back and show them the advice they have received. Advisees then pick whether to compete against high or low performers, with the outcome determined by their past performance on the quiz. Our key hypotheses are that (1) advisees who were anchored to high expectations would be more likely to receive advice to compete against the high performer group, and (2) that men are more likely than women to be advised to compete against the high-performer group when expectations are shown.

Methods

As part of our main experiment, participants choose to compete against a group of previous participants and have their expectations anchored based on their performance. Thus, we begin by first recruiting a sample of 50 participants from Prolific to complete a 10-question multiple choice mathematics quiz. The questions were taken from a paper-version of the ASVAB standardized exam, such that answers were not available online. Participants had five minutes to answer the quiz and were paid 10 cents for each correctly answered question. We define the top 20 scorers as the “High Performers” and the bottom 20 scorers as the “Low Performers.” On average, participants answered 4.42 questions correctly, High Performers scored between 5 and 10, and Low Performers scored between 0 and 3. In

order to anchor the expectations of participants in our main experiment, we simulated 1,000 pairings of groups of five participants, with the 5th percentile of groups scoring an average of 2.60 and the 95th percentile scoring an average of 6.40. We report these averages to participants in the Low Expectations and High Expectations treatment, respectively.

We then recruited 1,002 participants for the role of advisees in Stage 1 of our main experiment. To arrive at a gender-balanced sample, we dropped the last two male participants to complete the survey, ending up with a sample of 500 men and 500 women ($M_{\text{Age}} = 42.16$). Participants completed the same 10-item mathematics quiz as the earlier participants and were informed that their performance would affect their bonus earnings in a follow-up stage to be conducted a few days later. After completing the quiz, we informed them of the average score of a group of five participants from the preliminary survey. We randomly assigned them to learn about the 5th percentile of groups, which scored 2.60. (“Low Expectations” treatment) or the 95th percentile of groups, which scored 6.40 (“High Expectations” treatment). Participants then made a guess (unincentivized) about how many questions they think they answered correctly. The survey concluded with basic demographic questions.

We then recruited 1,000 participants for Stage 2, placing them in the role of advisers. We began by informing them of the mathematics quiz that participants in the preliminary study and Stage 1 had completed, and informed them of the average score of all participants in the preliminary study. Advisers had to recommend whether an advisee should compete against the Low Performers or the High Performers (we used these terms in the survey). We anticipated that being told to compete against High Performers is more flattering and hence being told to compete against the Low Performers would be disappointing if one had expected to do well. Advisees would earn 30 cents if their score was equal to or higher than a randomly selected member of the Low Performer group and they chose to compete against this group, or 50 cents if they performed as well or better than a randomly selected member

of the High Performer group and they chose to compete against that group. If they scored lower, they would not receive any bonus.

Advisers were randomly assigned to one of two treatments. In the “Baseline” treatment, they only observed the score of the advisee on the mathematics quiz. In the “Expectation” treatment, they observed the score as well as the advisee’s guess for how many questions they answered correctly. Note that since the outcome is determined only by the past score on the quiz, the advisee’s guess is immaterial to which group they should compete against. Moreover, in neither treatment did they receive any demographic information about their advisee. Participants gave recommendations to 10 advisees, which unbeknownst to them were five men and five women matched to have identical performance on the test.² They were informed that if their advice was shown to a participant who returned for the follow-up survey, they would receive the identical bonus as that participant. The survey then concluded with basic demographic questions.

Finally, we invited participants from Stage 1 back for the follow-up survey. Following our preregistration, we kept the survey open for 7 days. In total, 951 participants (481 men, 470 women) returned. The brief survey reminded them of the task they completed in Stage 1, informed them that other participants from Prolific had observed their real score and given them advice against which group to compete against, and finally were reminded them of how many questions they guessed they had answered correctly. Importantly, they were not informed of their true score or the score of the groups they could compete against. Participants then observed the advice from a randomly selected adviser and made their decision.

Results

We begin by examining the performance of the Stage 1 participants. Because our treatment takes place after participants completed the mathematics quiz, we would not

² We made this decision to account for the possibility of gender differences in performance.

expect a difference in performance across the Low and High Expectations treatments. Indeed, the two groups scored no different from one another (4.79 and 4.55 for the Low Expectations and High Expectations treatments, respectively, $t(998) = -1.64$, $p = .102$). The expectations treatment did, however, affect how well they thought they performed. Participants in the “Low Expectations” treatment guessed a score of 4.79 vs. 4.42 in the High Expectations treatment ($t(998) = 3.78$, $p < .001$), showing that the manipulation was successful, albeit small. Contrary to our expectations, we did find a gender difference in performance: men scored 5.02 on average, while women scored 4.33 ($t(998) = -4.73$, $p < .001$). Consistent with this difference, men thought they answered more questions correctly than did women (4.75 vs 3.56, $t(998) = -8.60$, $p < .001$). This difference, however, does not affect the interpretation of our findings, which will rely on an interaction of gender with an experimental treatment for Stage 2 participants.

We define a measure of “overconfidence” as (Performance - Estimate). We observe that women underestimate their performance by 0.76 points, while men do so by only 0.27 points $\Delta M = -0.49$, 95% CI $[-0.73, -0.25]$, $t(997.87) = -4.03$, $p < .001$. Although neither gender is overconfident, men on average are more confident in their performance than are women, as we had expected.

Next, we examine whether advisers took the advisee’s expectations into account. Column 1 of Table 3 reports a linear probability model on advising to compete against the High Performer group for advisers in the “Expectations” treatment, controlling for the true performance of the advisee. Because each adviser made ten recommendations, we cluster standard errors at the adviser level. However, contrary to our expectations, we do not see a significant effect of the expectations treatment. This may be because the induced difference in expectations was too small.

However, recall that men were more confident in their performance than were women. Our theory thus predicts that showing expectations should lead to more flattering advice

Table 3

Column 1 displays the advice given to compete against a High Performance group based on whether the advisees were primed with low or high expectations and the expectation level shown to the advisor. Column 2 displays the advice given to compete against a High Performance group based on the gender of the targets and the expectation level shown to the advisor. Column 3 displays the expected bonus of the advice received based on the gender of the targets and the expectation level shown to the advisor. All standard errors clustered at the level of the advisor.

	(1)	(2)	(3)
High Expectation	0.012 (0.011)		
Performance	0.128*** (0.003)		
Expectation Shown		-0.034* (0.013)	0.000 (0.003)
Advisee Male		0.002 (0.005)	0.000 (0.001)
Expectation x Male		0.032*** (0.009)	-0.005** (0.002)
Constant	-0.263*** (0.015)	0.339*** (0.010)	0.258*** (0.002)
N	5000	10 000	10 000
+ p < 0.1, * p < 0.05, ** p < 0.01, *** p < 0.001			

(i.e., a recommendation to compete against the High Performer group) for men than for women. We report a linear probability model with advice to compete against the high group as the outcome measure, and the advisee's gender, whether expectations were shown to the adviser, and the interaction of the two in Column 2 of Table 3. As predicted, we find a significant interaction effect: men are more likely to be advised to compete against the High Performers when expectations are shown than when they are hidden ($p < 0.001$). We show this result graphically in Figure 5).

To examine the quality of advice, we computed the expected bonus earnings for someone who followed the recommendations. For example, if an adviser suggested competing against the High Performer group, we matched the advisee against all 20 members of that group and determined how often their score matched or exceeded that of the member. We




Figure 5

Advisers observed real performance and the participants' estimated performance, but not their gender. As a result of their higher expectations, men (blue) were advised to compete against the high performance group more often than women (red). The grey line shows advice absent expectations, which did not differ by gender.

then multiplied this number by the respective bonus earnings (50 cents and 30 cents for the High and Low Performer group, respectively). To see if including expectations leads to worse advice for men, we report a linear probability model with the experimental treatment of the adviser, the gender of the advisee, and their interaction in Column 3 of Table 3. Displaying expectations led men to be advised to compete against the High Performer group more often, and this advice turned out to be bad: men receive worse advice than do women when expectations are displayed, but not in their absence.³

To determine whether flattering advice is truly costly, however, we need to examine the outcome of the advisees. In particular, they could ignore flattering advice, recognizing it as such and thus failing to adhere to it. Align with our prediction, participants who were

³ This analysis was not preregistered, and we note here that the reduction in expected earnings is small. However, it is interesting that expectations have a negative effect for men who underestimate their performance on average. One possibility is that advisers suggest the High Performer group more often than is optimal. We return to the possibility that advice is overall biased to be flattering in Study 2.

Table 4

Column 1 displays the actual bonus based on the performance and whether the participant is primed with high expectations. Column 2 displays the actual bonus based on the performance and gender of the participant. Columns 3 and 4 display the chance of adopting the advice based on the gender of the participants, whether the advisee sees expectations, and whether the advice is to compete against high performers. The former considers only the main effect, while the latter also includes the interactive effect.

	(1)	(2)	(3)	(4)
Performance	0.038*** (0.003)	0.039*** (0.003)		
High Expectation	-0.020 (0.013)			
Advisee Male		-0.002 (0.013)	-0.045+ (0.026)	-0.144*** (0.041)
Expectation x Male				0.113* (0.053)
Expectation Shown			-0.002 (0.026)	-0.057 (0.040)
Advice: High Performer			-0.138*** (0.028)	-0.207*** (0.051)
Advice: High Performer x Male				0.131* (0.056)
Expectation x Advice: High Performer				-0.009 (0.056)
Constant	0.087*** (0.016)	0.076*** (0.015)	0.857*** (0.024)	0.906*** (0.030)
N	483	483	951	951

+ $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

primed with high expectations earned less when their expectations were conveyed to the advisers, although this result is only directional (see Column 1 of Table 4). Similarly, as shown in Column 2 of Table 4, we also found that male participants received less reward when their expectations were presented, even when their performance was the same. These findings suggest that flattering advice is not without consequences. Furthermore, our research indicates that because males are more likely to follow such advice (see Column 4 of Table 4, their tendency towards competition intensifies the cost of such advice.

Discussion

men underestimate their score on a mathematics test less than do women. When these expectations are displayed to advisers, they are more likely to tell men to compete against the High Performer group. Notably, this is bad advice: men whose advisers observed their expectations got worse advice as a result. Our findings suggest that males with the same scores as their female counterparts earned less, although the evidence is only directional. This may be due to males receiving more favorable advice and being more likely to adopt it, resulting in even worse outcomes.

General Discussion

Advice has the potential to shape people’s career and personal outcomes. Honest feedback, however, may be painful to learn if it falls short of one’s expectations. As prior work notes, this may motivate people to avoid information and avoid seeking help (Bénabou et al., 2022; Golman et al., 2017; Jaroszewicz et al., 2021). We present evidence from two experiments that advisers are also cognizant of this cost. As a result, they present flattering advice that avoids disappointing the recipient, and correctly anticipate that this boosts how advisees perceive them. However, this flattering advice comes at a cost to advisees, who would do worse if they followed it blindly, as we show in Study 1.

Moreover, a desire to avoid disappointment also means that advisers have to take into account the expectations of the advisee. We document that men are more optimistic about their performance than women are. As a result, they receive more flattering advice and are more likely to be told to aim higher. Notably, in the context of our experiment, this turns out to be bad advice ex post.

Our findings have implications for organizational practice, where mentoring and advice-giving may take into account an employees’ expectation. We document this as a novel source of gender bias. Organizations could reduce this bias by calibrating employees’ expectations to reduce overconfidence.

Participants in our experiment were paired anonymously. Even so, we document this desire to avoid disappointment. We anticipate that advice would be more flattering in face-to-face communication and without anonymity. Moreover, existing relationships might make it even more difficult for advisers to be honest.

In our experiments, participants only received advice once and did not evaluate the adviser after observing the outcome of their decision. For example, advice that leads to bad outcomes may undermine the interpersonal benefits of flattery. Alternatively, people may still like the flattering advice and not fault the adviser for the bad outcome. Moreover, our setting involved only a single piece of advice on one task. Future research could examine whether people return to those who gave them flattering advice, or if they prefer someone who gave them the honest (but unpleasant) truth.

Advice has long been studied from the perspective of the receiver. Similarly, research on belief utility has examined how recipients respond to the valence of the information they receive. Here, we show that advisers, too, take into account the psychological impact of the information they convey. They may have even greater motivations to avoid conveying bad news, because they incur the interpersonal costs of delivering unfavorable information without reaping the benefits from helping someone make a better choice.

References

- Ache, F., Rader, C., & Hütter, M. (2020). Advisors want their advice to be used—but not too much: An interpersonal perspective on advice taking. *Journal of Experimental Social Psychology*, 89, 103979.
- Bénabou, R., Jaroszewicz, A., & Loewenstein, G. (2022). *It Hurts to Ask*. National Bureau of Economic Research.
- Benjamin, D., & Budescu, D. V. (2015). Advice from experience: Communicating incomplete information incompletely. *Journal of Behavioral Decision Making*, 28(1), 36–49.
- Exley, C., & Nielsen, K. (2022). The gender gap in confidence: Expected but not accounted for. *Available at SSRN 4352381*.
- Fiske, S. T., Cuddy, A. J., & Glick, P. (2007). Universal dimensions of social cognition: Warmth and competence. *Trends in Cognitive Sciences*, 11(2), 77–83.

- Gill, D., & Prowse, V. (2012). A structural analysis of disappointment aversion in a real effort competition. *American Economic Review*, 102(1), 469–503.
- Golman, R., Hagmann, D., & Loewenstein, G. (2017). Information avoidance. *Journal of Economic Literature*, 55(1), 96–135. <https://doi.org/10.1257/jel.20151245>
- Jaroszewicz, A., Loewenstein, G., & Bénabou, R. (2021). *The pain of asking and being asked for informal help*.
- John, L. K., Jeong, M., Gino, F., & Huang, L. (2019). The self-presentational consequences of upholding one’s stance in spite of the evidence. *Organizational Behavior and Human Decision Processes*, 154, 1–14.
- Jonas, E., & Frey, D. (2003). Information search and presentation in advisor–client interactions. *Organizational Behavior and Human Decision Processes*, 91(2), 154–168.
- Kanze, D., Huang, L., Conley, M. A., & Higgins, E. T. (2018). We ask men to win and women not to lose: Closing the gender gap in startup funding. *Academy of Management Journal*, 61(2), 586–614.
- Kőszegi, B., & Rabin, M. (2006). A model of reference-dependent preferences. *The Quarterly Journal of Economics*, 121(4), 1133–1165.
- Kray, L. J. (2000). Contingent weighting in self-other decision making. *Organizational Behavior and Human Decision Processes*, 83(1), 82–106.
- Kray, L., & Gonzalez, R. (1999). Differential weighting in choice versus advice: I’ll do this, you do that. *Journal of Behavioral Decision Making*, 12(3), 207–218.
- Loewenstein, G. (1987). Anticipation and the valuation of delayed consumption. *The Economic Journal*, 97(387), 666–684.
- Loewenstein, G., & Molnar, A. (2018). The renaissance of belief-based utility in economics. *Nature Human Behaviour*, 2(3), 166–167.
- Shalvi, S., Soraperra, I., Weele, J. J. van der, & Villeval, M. C. (2019). *Shooting the messenger? Supply and demand in markets for willful ignorance*.